

```

/* =====
 *
 * CS 566 - Assignment 04
 * Camillo Lugaresi, Cosmin Stroe
 *
 * This code solves the Travelling Salesman Problem by using a
 * Depth First Search, Branch and Bound algorithm.
 *
 * ===== */

#include <stdio.h>
#include <limits.h>
#include <stddef.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include "mpi.h"
#include "tsplib95.h"
#include "tsp.h"

struct tsp_state *tsp_state_alloc(struct tsp_matrix *matrix)
{
    struct tsp_state *state;
    state =
        (struct tsp_state *)malloc(sizeof(struct tsp_state, tour) +
                                   sizeof(state->tour[0]) * matrix->n);

    state->cost = 0;
    state->ub = INT_MAX;
    state->ub_rank = 0;
    state->give_depth = state->subtree_depth = 1;
    state->len = 1;
    state->tour[0] = 0; /* we always begin the path from node 0 */
    /* state->first_given = state->last_given = NULL; */
    state->matrix = matrix;
    state->last_started = calloc(matrix->n, sizeof(*state->last_started));
    state->term_token.request = MPI_REQUEST_NULL;
    state->term_token.token = NO_TOKEN;
    state->term_token.mycolor = WHITE;
    state->work_partner = 0;
    state->best_tour = calloc(matrix->n, sizeof(*state->best_tour));
    return state;
}

int main(int argc, char *argv[])
{
    struct tsp_matrix matrix;
    struct tsp_state *state;
    int numprocs, namelen, i;
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    FILE *file;

    /* setup */
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    MPI_Get_processor_name(processor_name, &namelen);

    /* parse our TSP matrix */
    file = fopen(argv[1], "r");
    parse_matrix_from_file(&matrix, file);

    /* setup the state variable */
    state = tsp_state_alloc(&matrix);
    MPI_Comm_rank(MPI_COMM_WORLD, &(state->myrank));
    MPI_Comm_size(MPI_COMM_WORLD, &(state->num_procs));

    /* printf("%d: After tsp_state_alloc\n", state->myrank); */

```

```

    if (state->myrank == 0) {
        state->work_state = WORKING;
        print_matrix(&matrix);
    } else {
        state->work_state = NEED_WORK;
    }

    tsp(state);

    if (state->myrank == 0) {
        fprintf(stdout, "best: %d\n", state->ub);
        for (i = 0; i < state->matrix->n; i++)
            fprintf(stdout, "%2d ", state->best_tour[i]);
        fprintf(stdout, "\n");
        fprintf(stdout, "total time: %f\n", state->total_time);
        fprintf(stdout, "work time: %f (%2.1f%%)\n", state->work_time,
            state->work_time / state->total_time * 100.0);
    }

    MPI_Finalize();
    return 0;
}

void tsp(struct tsp_state *state)
{
    MPI_Status pending_status;
    MPI_Request request;
    int i;

    double start_time = MPI_Wtime();
    state->work_time = 0;

    /* begin by giving away work */
    if (state->myrank == 0) {
        service_pending_messages(state);
    }

    while (state->work_state != QUIT) {

        if (state->work_state == NEED_WORK) {
            request_work(state);

            if (state->work_state != WORKING) {
                state->work_state = IDLE;
                if (state->myrank == 0
                    || state->term_token.token != NO_TOKEN)
                    send_token(state);
            }
        }

        if (state->work_state == WORKING) {
            double work_start_time = MPI_Wtime();
            do_work(state);
            state->work_time += MPI_Wtime() - work_start_time;
        }

        if (state->work_state == IDLE) {
            fprintf(stdout, "%d: Probing for pending messages.\n", state->myrank);
            MPI_Probe(MPI_ANY_SOURCE, MPI_ANY_TAG, MPI_COMM_WORLD,
                &pending_status);
            fprintf(stdout, "%d: We have messages.\n", );
        }

        service_pending_messages(state);
    }
}

```

```

    if (state->myrank == 0) {
        if (state->ub_rank != state->myrank) {
            MPI_Isend(&state->myrank, 1, MPI_INT, state->ub_rank,
                     BEST_PATH_REQ_TAG, MPI_COMM_WORLD, &request);
            MPI_Recv(state->best_tour, state->matrix->n, MPI_INT, state->ub_rank,
                     BEST_PATH_TAG, MPI_COMM_WORLD, &pending_status); // send the termination message
        }
        for (i = 1; i < state->num_procs; i++) {
            MPI_Isend(&i, 1, MPI_INT, i, TERMINATION_TAG,
                     MPI_COMM_WORLD, &request);
        }
    }

    state->total_time = MPI_Wtime() - start_time;
}

void send_token(struct tsp_state *state)
{
    MPI_Request request;
    if (state->term_token.token == NO_TOKEN) {
        state->term_token.token = WHITE;
    }
    MPI_Isend(&state->term_token.token, 1, MPI_INT,
              ((state->myrank + 1) % state->num_procs), TOKEN_TAG,
              MPI_COMM_WORLD, &request);
    state->term_token.mycolor = WHITE;
}

int next_available_node(struct tsp_state *state, int depth)
{
    int next_node = state->last_started[depth] + 1, last_next_node, i;
    do {
        last_next_node = next_node;
        for (i = 0; i < depth; i++)
            if (next_node == state->tour[i])
                next_node++;
    } while (last_next_node != next_node);
    return next_node;
}

/*
    do a fixed amount of work
*/
void do_work(struct tsp_state *state)
{
    int work_counter = 0, i, next_node;
    int progress_counter = 0;

    for (work_counter = 0; work_counter < WORK_SLICE; work_counter++) {
        int go_up = 1;

#ifdef 0
        if (progress_counter == 5000000) {
            fprintf(stdout, "%d: ", state->myrank);
            for (i = 0; i < state->len; i++)
                fprintf(stdout, "%2d ", state->tour[i]);
            fprintf(stdout, "\n");
            progress_counter = 0;
        } else
            progress_counter++;

#endif

        /* try to go down */
        if (state->len < (state->matrix->n)) {
            /* pick the next available node */
            next_node = next_available_node(state, state->len);

```

```

        //fprintf(stdout, "next_node: %d\n", next_node);

        if (next_node < state->matrix->n) {

            /* go down */
            state->last_started[state->len] = next_node;

            /* append node to tour */
            state->tour[state->len] = next_node;
            state->cost +=
                CELL(state->matrix,
                    state->tour[state->len - 1],
                    state->tour[state->len]);
            state->len++;

            //fprintf(stdout, "down %d, len %d\n", next_node, state->len);

            /* reset next level */
            if (state->len < state->matrix->n)
                state->last_started[state->len] = 0;

            if (state->cost >= state->ub) {
                /* prune bad branches */
                fprintf(stdout, "pruning: ");
                for (i = 0; i < state->len; i++) fprintf(stdout, "%2d ", state-
>tour[i]);

                fprintf(stdout, "\n");*/
                go_up = 1;
            } else {
                go_up = 0;

                /* if we have a complete tour, update the ub */
                if (state->len == state->matrix->n) {
                    //for (i = 0; i < state->len; i++) printf("%2d ", state-
>tour[i]);

                    int cycle_cost =
                        state->cost +
                        CELL(state->matrix,
                            state->tour[state->
                                len - 1],
                            state->tour[0]);
                    //printf("cost: %d", cycle_cost);
                    if (cycle_cost < state->ub) {
                        state->ub = cycle_cost;
                        state->ub_rank =
                            state->myrank;
                        send_ub_message(state);
                        //printf(" new best");
                        memcpy(state->best_tour,
                            state->tour,
                            sizeof(*state->
                                tour) *
                                state->len);
                    }
                    //printf("\n");
                }
            }
        }
    }

    if (go_up) {
        printf("up\n");
        /* go up */
        state->len--;
        if (state->len < state->subtree_depth) {
            state->work_state = NEED_WORK;
            break;
        }
    }
}

```

```

        }
        state->cost -=
            CELL(state->matrix, state->tour[state->len - 1],
                state->tour[state->len]);
    }
}

/*
 * service all the pending messages in our message queue
 */
void service_pending_messages(struct tsp_state *state)
{
    MPI_Status pending_status, status;
    MPI_Request request;
    int msg_pending;
    /* is there a message pending in the message QUEUE ? */
    int temp, possible_ub;

    /* service UB broadcasts */
    MPI_Iprobe(MPI_ANY_SOURCE, UB_TAG, MPI_COMM_WORLD, &msg_pending,
        &pending_status);
    while (msg_pending) {
        MPI_Recv(&possible_ub, 1, MPI_INT, pending_status.MPI_SOURCE,
            pending_status.MPI_TAG, MPI_COMM_WORLD, &status);
        if (possible_ub < state->ub) {
            state->ub = possible_ub;
            state->ub_rank = pending_status.MPI_SOURCE;
        }
        MPI_Iprobe(MPI_ANY_SOURCE, UB_TAG, MPI_COMM_WORLD, &msg_pending,
            &pending_status);
    }

    /* service BEST PATH REQ message */
    MPI_Iprobe(MPI_ANY_SOURCE, BEST_PATH_REQ_TAG, MPI_COMM_WORLD,
        &msg_pending, &pending_status);
    while (msg_pending) {
        MPI_Recv(&temp, 1, MPI_INT, pending_status.MPI_SOURCE,
            pending_status.MPI_TAG, MPI_COMM_WORLD, &status);

        MPI_Isend(state->best_tour, state->matrix->n, MPI_INT,
            pending_status.MPI_SOURCE, BEST_PATH_TAG,
            MPI_COMM_WORLD, &request);

        MPI_Iprobe(MPI_ANY_SOURCE, BEST_PATH_REQ_TAG, MPI_COMM_WORLD,
            &msg_pending, &pending_status);
    }

    /* service work requests */
    MPI_Iprobe(MPI_ANY_SOURCE, WORK_REQ_TAG, MPI_COMM_WORLD, &msg_pending,
        &pending_status);
    while (msg_pending) {
        MPI_Recv(&temp, 1, MPI_INT, pending_status.MPI_SOURCE,
            pending_status.MPI_TAG, MPI_COMM_WORLD, &status);
        service_work_request(state, pending_status);
        MPI_Iprobe(MPI_ANY_SOURCE, WORK_REQ_TAG, MPI_COMM_WORLD,
            &msg_pending, &pending_status);
    }

    /* service token sends */
    MPI_Iprobe(MPI_ANY_SOURCE, TOKEN_TAG, MPI_COMM_WORLD, &msg_pending,
        &pending_status);
    while (msg_pending) {
        /* fprintf(stdout, "%d: Receiving token from %d.\n", state->myrank,
pending_status.MPI_SOURCE); */
    }
}

```

```

        MPI_Recv(&(state->term_token.token), 1, MPI_INT,
                pending_status.MPI_SOURCE, pending_status.MPI_TAG,
                MPI_COMM_WORLD, &status);
        /* fprintf(stdout, "%d: Token color: %d.\n", state->myrank, state->term_token.token); */
        if (state->term_token.mycolor == BLACK)
            state->term_token.token = BLACK;
        if (state->myrank == 0) {
            if (state->term_token.token == WHITE)
                state->work_state = QUIT;
            else
                state->term_token.token = WHITE;
        }
        if (state->work_state == IDLE) {
            send_token(state);
        }
        MPI_Iprobe(MPI_ANY_SOURCE, TOKEN_TAG, MPI_COMM_WORLD,
                &msg_pending, &pending_status);
    }

    /* service TERMINATION broadcast */
    MPI_Iprobe(MPI_ANY_SOURCE, TERMINATION_TAG, MPI_COMM_WORLD,
                &msg_pending, &pending_status);
    while (msg_pending) {
        state->work_state = QUIT;
        return;
    }
}

void service_work_request(struct tsp_state *state, MPI_Status status)
{
    int work_deny = 1;
    int i;
    int *outbuf;
    int next_node;

    if (state->work_state == WORKING) {
        while (state->give_depth < (state->matrix->n - MIN_WORK_LEVELS)
                && (next_node =
                    next_available_node(state,
                                        state->give_depth)) >=
                    state->matrix->n) {
            state->give_depth++;
        }

        /* do we have work to give away ? */
        if (state->give_depth < (state->matrix->n - MIN_WORK_LEVELS)
            && next_node < state->matrix->n) {
            /* ok, give out the work */
            work_deny = 0;
            state->last_started[state->give_depth] = next_node;
            /* give away the next node at this depth */

            /* prepare the message buffer(we 're sending the current path to the given depth,
with the new last node) */
            outbuf =
                calloc(state->give_depth + 1,
                    sizeof(state->tour[0]));
            for (i = 0; i < state->give_depth; i++) {
                outbuf[i] = state->tour[i];
            }
            outbuf[state->give_depth] =
                state->last_started[state->give_depth];

            /* synchronous send, because the receiver should be already waiting for the reply
(and we want to free the memory) */
            fprintf(stdout, "%d: Giving work to %d:", state->myrank,

```

```

        status.MPI_SOURCE);
    for (i = 0; i < state->give_depth + 1; i++) {
        fprintf(stdout, " %2d", outbuf[i]);
    }
    fprintf(stdout, "\n");

    MPI_Send(outbuf, state->give_depth + 1, MPI_INT,
             status.MPI_SOURCE, WORK_ACK_TAG,
             MPI_COMM_WORLD);

    free(outbuf);

    /* update our token */
    if (state->myrank > status.MPI_SOURCE)
        state->term_token.mycolor = BLACK;
}

if (work_deny) {
    /* no, we don't have any work to give away */
    fprintf(stdout,
            "%d: Denying work request from %d (ws:%d, gd:%d, nn:%d)\n",
            state->myrank, status.MPI_SOURCE, state->work_state,
            state->give_depth, next_node);
    MPI_Send(&work_deny, 1, MPI_INT, status.MPI_SOURCE,
            WORK_ACK_TAG, MPI_COMM_WORLD);
}
}

void request_work(struct tsp_state *state)
{
    int got_work = 0;
    int msg_size, i;
    MPI_Status status;
    MPI_Request req;
    int initial_partner = state->work_partner;
    int reply_available;

    do {
        if (state->work_partner != state->myrank) {
            printf("%d: Sending work request to %d\n", state->myrank, state->work_partner);
            MPI_Isend(&msg_size, 1, MPI_INT, state->work_partner,
                    WORK_REQ_TAG, MPI_COMM_WORLD, &req);

            reply_available = 0;
            while (!reply_available) {
                MPI_Probe(MPI_ANY_SOURCE, MPI_ANY_TAG,
                        MPI_COMM_WORLD, &status);
                service_pending_messages(state);

                /* check for request reply */
                MPI_Iprobe(state->work_partner, WORK_ACK_TAG,
                        MPI_COMM_WORLD, &reply_available,
                        &status);
            }

            MPI_Get_count(&status, MPI_INT, &msg_size);

            MPI_Recv(&state->tour, msg_size, MPI_INT,
                    state->work_partner, WORK_ACK_TAG,
                    MPI_COMM_WORLD, &status);
            if (msg_size > 1) {
                printf("%d: Got work from %d\n", state->myrank, state->work_partner);
                state->work_state = WORKING;
                state->give_depth = state->subtree_depth =
                    msg_size;
                state->len = msg_size;
            }
        }
    } while (!got_work);
}

```

```
        state->last_started[state->give_depth] = 0;

        state->cost = 0;
        for (i = 1; i < state->len; i++) {
            state->cost +=
                CELL(state->matrix,
                    state->tour[i - 1],
                    state->tour[i]);
        }

        got_work = 1;
    }
}

state->work_partner =
    (state->work_partner + 1) % state->num_procs;
}
while (state->work_partner != initial_partner && !got_work);
}

void send_ub_message(struct tsp_state *state)
{
    int i;
    MPI_Request request;

    for (i = 0; i < state->num_procs; i++) {
        if (i == state->myrank)
            continue;
        MPI_Isend(&state->ub, 1, MPI_INT, i, UB_TAG, MPI_COMM_WORLD,
            &request);
    }
}
```